

# Week 5: Neural Nets & SVM



<https://www.rambus.com/blogs/fpgas-take-on-convolutional-neural-networks/>

# Week 4 Review: Trees

- What is the criterion for splitting classification?  
Regression?
  -
- What is a 1-layer tree called?
  -
- What is different about random forests compared with a single decision tree?
  -
- What are boosting methods fitting after the first tree?
  -

# Week 4 Review: Trees

- What is the criterion for splitting classification? Regression?
  - Classification – Gini or Entropy (info. Gain)
  - Regression –  $R^2$  or another SSE-like metric (MAE, MSE, etc)
- What is a 1-layer tree called?
  - A stump
- What is different about random forests compared with a single decision tree?
  - Random forests use many decision trees, using bootstrapping of samples and randomly sampling the features for each tree (mistake here, what is it?)
- What are boosting methods fitting after the first tree?
  - The residuals (  $(\text{predictions} - \text{actual})^2$  )

# Week 4 review quiz


- Using the `auto.dt.clean.csv` dataset (under week 5), predict mpg from everything else with a random forest
- Create train/test sets
- Try 3 values for `mtry`, including the default
- Evaluate performance on the train and test set, and check for overfitting
- Plot the feature importances and explain them

# SVM

- Support vector machine
- Invented in 1963 by Vladimir N. Vapnik and Alexey Ya. Chervonenkis
  - Vlad and 2 others intro'd the 'kernel trick' in 1992
- Works well for fewer data points (<20k) and large number of features


# How SVMs work

- Mathematically it finds the optimal separating hyperplane between classes
- A hyperplane is a dot in 1D, a line in 2D, a plane in 3D, and a hyperplane in 4D+




# How SVMs work

- Finds the optimal by finding hyperplane with max distance between the 2 nearest points (called support vectors)



# Overlapping points

- If classes are overlapping, we can penalize the misclassification. The penalty is proportional to  $C$ , so a large  $C$  will avoid misclassification




$C$  is a tunable hyperparameter




# Kernel Trick

- We can transform the data with a kernel so it is separable




# Kernel Trick

- Most common kernels are polynomial (degree of polynomial is a hyperparameter) and radial basis function (RBF, Gaussian)
- These essentially map data from one space to another (it's a bit more complicated than that, but gets into complex math)



# Kernel Trick



$$K(x, x') = \exp\left(-\frac{(\|x - x'\|)^2}{2\sigma}\right)$$

Sigma is a tunable hyperparameter

[https://www.youtube.com/watch?v=H\\_I0pYdzBSk](https://www.youtube.com/watch?v=H_I0pYdzBSk)

<http://web.mit.edu/6.034/wwwbob/svm-notes-long-08.pdf>

# SVM Runtime


- The runtime depends on the hyperparameter  $C$
- $O(\max(n,d) \min(n,d)^2)$ ,  $n$  = number of samples,  $d$  = number of features
- In general,  $n^2*d$  for RBF, and  $n*d$  for linear kernels
- Low  $C$  (small misclassification penalty) usually decreases runtime

Chapelle, Olivier. "Training a support vector machine in the primal." *Neural Computation* 19.5 (2007): 1155-1178.

# More SVM Resources

- The SVM algorithm has a *lot* more math behind it, here are some follow-up resources
  - <http://blog.hackerearth.com/simple-tutorial-svm-parameter-tuning-python-r>
  - <http://www.robots.ox.ac.uk/~cvrg/bennett00duality.pdf>
  - <https://www.svm-tutorial.com/>
  - <https://jakevdp.github.io/PythonDataScienceHandbook/05.07-support-vector-machines.html>
  - [http://scikit-learn.org/stable/auto\\_examples/svm/plot\\_rbf\\_parameters.html](http://scikit-learn.org/stable/auto_examples/svm/plot_rbf_parameters.html)
  - [https://github.com/eriklindernoren/ML-From-Scratch/blob/master/mlfromscratch/supervised\\_learning/support\\_vector\\_machine.py](https://github.com/eriklindernoren/ML-From-Scratch/blob/master/mlfromscratch/supervised_learning/support_vector_machine.py)

- Support vector machines (SVMs) can also be used for regression (SVR)
- Similar (complicated) math



<https://www.mathworks.com/help/stats/understanding-support-vector-machine-regression.html>

<http://www.svms.org/regression/SmSc98.pdf>

# SVM demo/exercise in R

- On heart disease, using caret, kernlab, and e1071 packages

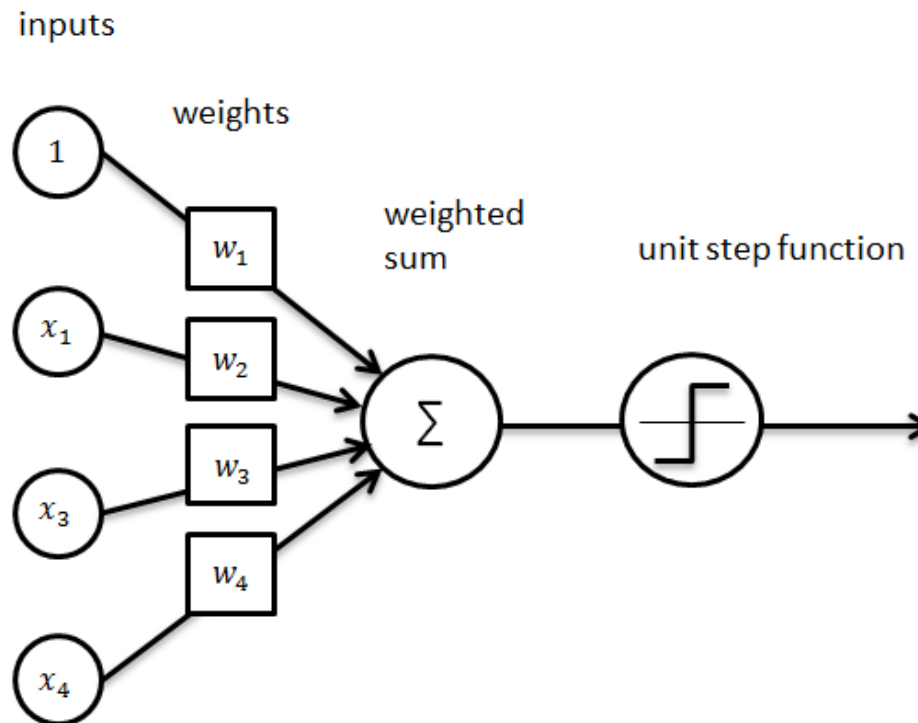
# Neural networks

- Can capture most any statistical pattern due to highly nonlinear behavior
- Infinite possibilities for configurations (architecture/topology)
- Requires GPUs (now Google's TPU as well) to train big nets
- Outperforms all other algorithms when set up correctly
- Difficult to understand what the net is learning/doing




# Perceptron

- If the weighted sum is greater than 0, output 1, otherwise 0





# Perceptron

- If the weighted sum is greater than 0, output 1, otherwise 0



# Activation functions and bias

- Can add a bias to shift the output left or right on the x-axis





<https://stackoverflow.com/a/2499936/4549682>

<https://www.neuraldesigner.com/blog/perceptron-the-main-component-of-neural-networks>

# Activation functions and bias

- Can add a bias to shift the output left or right on the x-axis





<http://www.snee.com/bobdc.blog/2017/09/understanding-activation-funct.html>

[https://en.wikipedia.org/wiki/Activation\\_function#Comparison\\_of\\_activation\\_functions](https://en.wikipedia.org/wiki/Activation_function#Comparison_of_activation_functions)

<https://www.neuraldesigner.com/blog/perceptron-the-main-component-of-neural-networks>

# Activation functions and bias

- ReLU (rectified linear unit) almost always used, although ELU sometimes works better




<https://arxiv.org/abs/1511.07289>

<http://laid.delanover.com/activation-functions-in-deep-learning-sigmoid-relu-lrelu-prelu-rrelu-elu-softmax/>

<https://www.neuraldesigner.com/blog/perceptron-the-main-component-of-neural-networks>

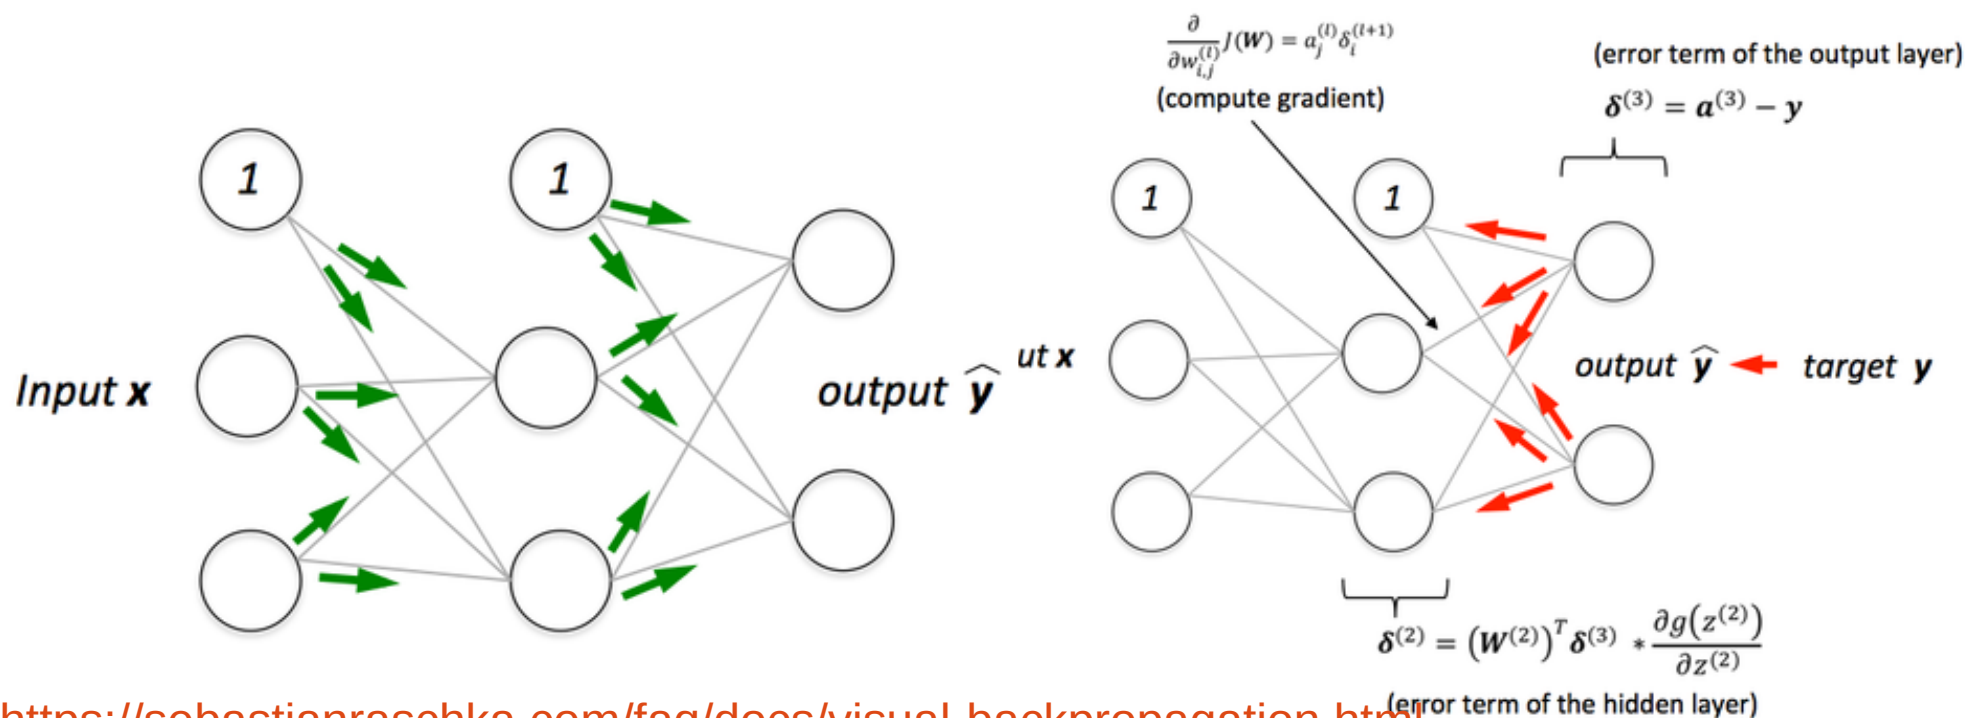
# Matrix math – dense layers

- Now we start stacking neurons on top of each other
- When we have a lot of hidden layers, the network becomes ‘deep’ and we have “deep learning”




# Loss functions and Backpropagation

- We need a way to find the best weights in the network
- Make a prediction, calculate error, and move the weights in the direction that minimizes error



# Loss functions and Backpropagation

- Keep doing this until our weights stop moving much
- Called gradient descent, because we are descending down the slope of the loss function



Link with the math:

<http://alexminnaar.com/deep-learning-basics-neural-networks-backpropagation-and-stochastic-gradient-descent.html>

<https://sebastianraschka.com/faq/docs/visual-backpropagation.html>



# Loss functions

- For regression, we can use MSE, MAE, or any other custom equation
- For classification, often categorical cross entropy (multiple classes) or binary cross entropy (2 classes)

<https://keras.io/losses/>

[http://ml-cheatsheet.readthedocs.io/en/latest/loss\\_functions.html](http://ml-cheatsheet.readthedocs.io/en/latest/loss_functions.html)

[https://en.wikipedia.org/wiki/Cross\\_entropy](https://en.wikipedia.org/wiki/Cross_entropy)

# Convolutional and pooling layers


- Used mostly for image recognition, but pattern recognition in general (i.e. for 1D time series)

1	1	1	0	0
0	1	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>
0	0	1 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>
0	0	1 <sub>x1</sub>	1 <sub>x0</sub>	0 <sub>x1</sub>
0	1	1	0	0

Image

4	3	4
2	4	3

Convolved  
Feature



Convolved  
feature

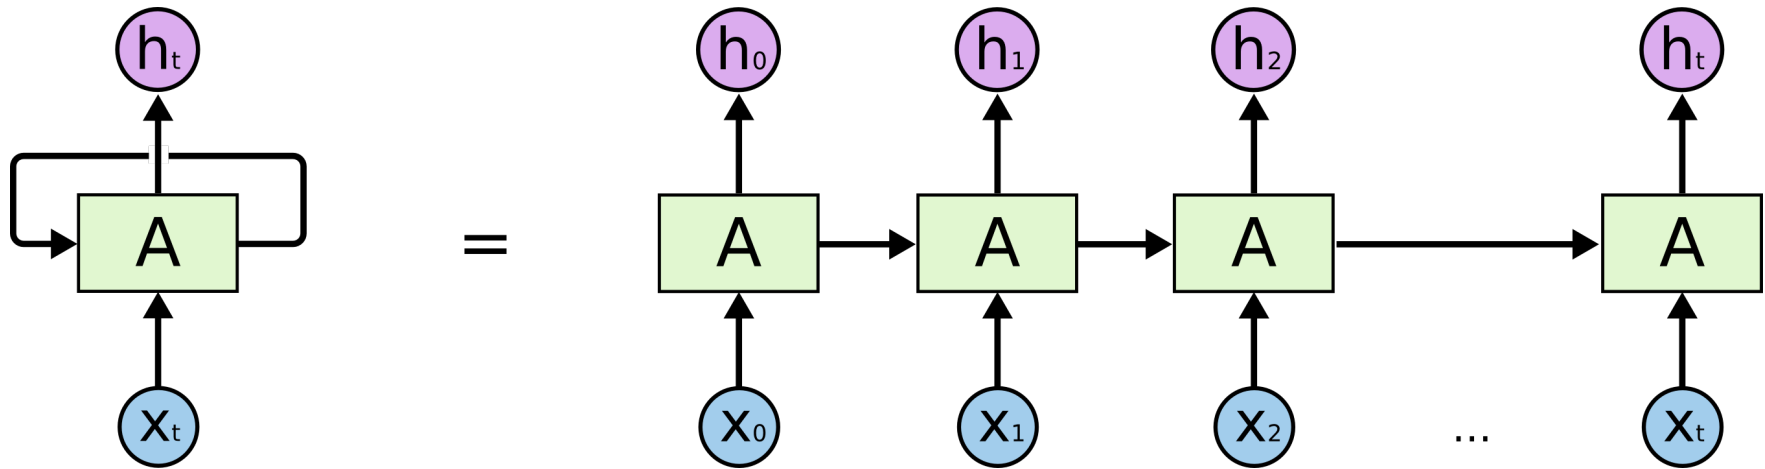
1	7
5	9

Pooled  
feature


<http://ufldl.stanford.edu/tutorial/supervised/FeatureExtractionUsingConvolution/>

<http://ufldl.stanford.edu/tutorial/supervised/Pooling/>

# Recurrent networks



LSTM cell (long short term memory):



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

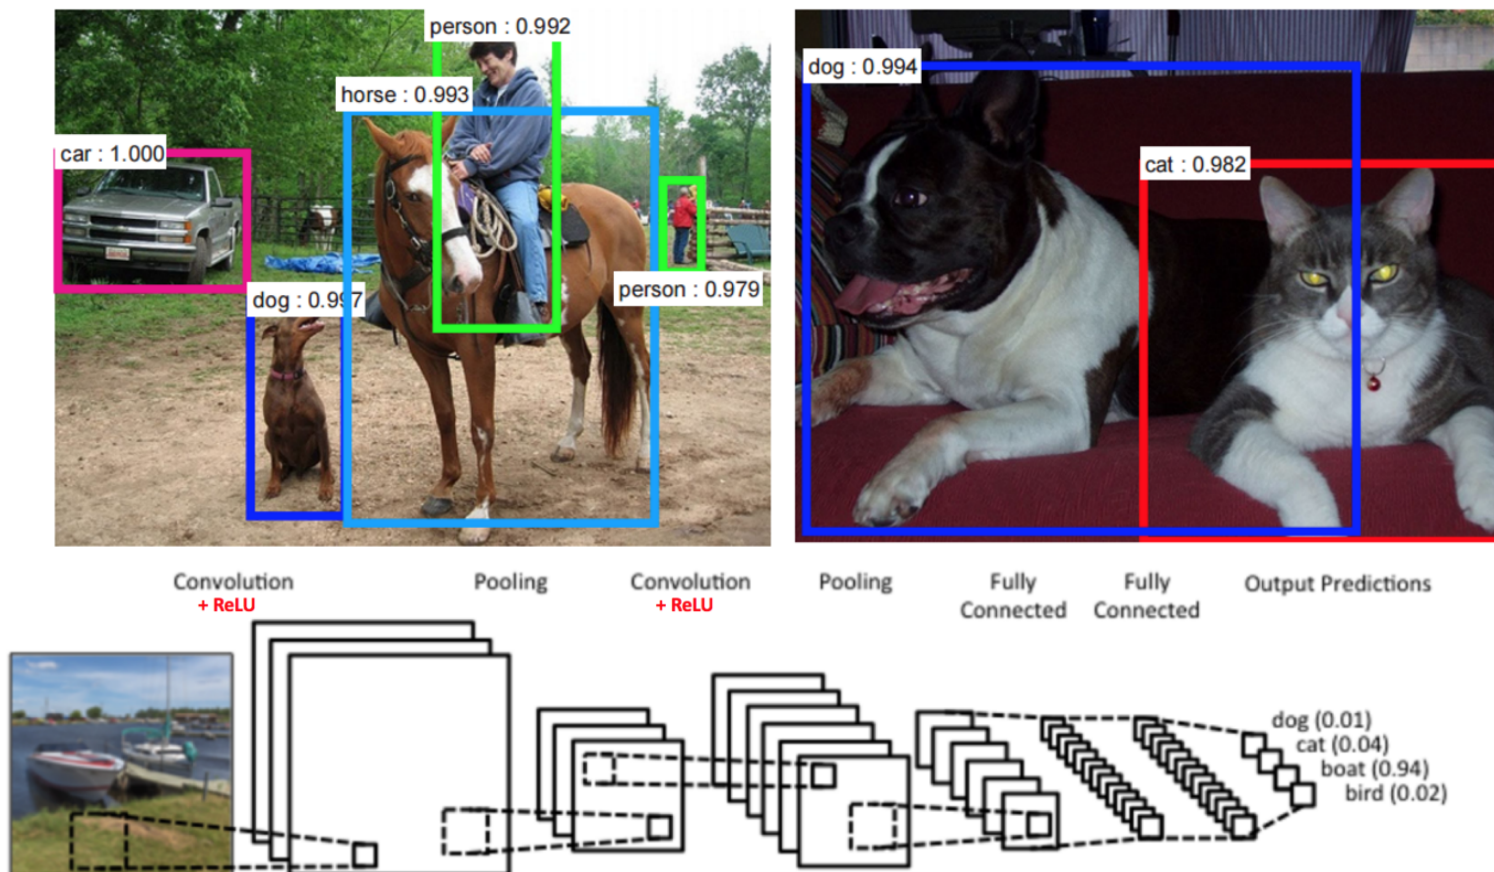
$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

# Image recognition

- These networks are very large and expensive to train



<https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

# Other applications

- Time series analysis/prediction
  - Finance (predicting future trends)
  - Audio (voice recognition/speech generation, **generating** /recommending music, recognizing songs like Shazam)
  - Electroencephelography (EEG) (brainwaves – predict if you want to move your arm, words you are imagining, etc)
  - Business/commerce data (predicting customer trends or outcomes)
- Generative adversarial networks (GANs)
- Speech/writing generation (RNN/LSTMs)
- Deep Q-learning (reinforcement learning with neural nets)

# Interesting applications

- Making art (**deepart**, **deepdream**)
- Google's speech generation improved by neural nets (**wavenet**)
- Deep reinforcement learning **beats Go** world champ
- **Deep reinforcement learning** playing video games, also going for **general AI**
- **Predicting what you are seeing** from your brainwaves
- **Face recognition**
- Creating neural nets...**with neural nets**

# Other neural network topics we didn't cover

- Embedding layers (and word2vec, etc)
- Backpropagation for RNNs
- Residual nets (e.g., ResNet50)
- Transfer learning
- Regularization
- Early stopping
- Optimizers
- Other neural net libraries ([lasagna](#), [torch](#), theano, CNTK...)

# Example/project: bike share neural net

- Data from here:  
<https://archive.ics.uci.edu/ml/datasets/bike+sharing+dataset>
- Train a regression neural net to predict the amount of bike rentals each hour
- Try at least 5 different architectures, pick the best one and defend your choice
- Post to the week 5 bike share exercise discussion